


## 1. Multi-threaded Monte Carlo Integration

Monte Carlo Integration is a numerical integration method. It can be used for determining the surface of a circle with radius  $r$ . This circle fits into a square of size  $2 \cdot r$  as shown in the following figure:



The surface of the square is known to be  $S_{\text{square}} = 4 \cdot r^2$ .

Assume we randomly choose  $N_{\text{try}}$  tuples  $(x, y)$  where  $-r \leq x, y \leq +r$ . Let  $N_{\text{hit}}$  be the number tuples with  $x^2 + y^2 \leq r^2$ . (In the figure above, these tuples are shown as red bullets.) Using the Monte Carlo Integration method we can estimate the surface of the circle as follows:

$$S_{\text{circle}} \simeq \frac{N_{\text{hit}}}{N_{\text{try}}} S_{\text{square}}. \quad (1)$$

You can compare this result with the known result:  $S_{\text{circle}} = \pi r^2$

Your task is to implement a multi-threaded program for computing  $S_{\text{circle}}$  for  $r = 1$  with the following properties:

- Use POSIX threads
- Allow for the number of threads to be  $N_{\text{thrd}} > 1$
- Use a Pthreads mutex for the concurrent update of  $N_{\text{try}}$  and  $N_{\text{hit}}$
- Let the worker threads run forever while using the main thread to periodically print the current estimate for  $S_{\text{circle}}$

Note that you need a pseudo-random number generator (RNG) that can be used in the context of multi-threading. `drand48_r` is a suitable choice as it allows to keep the RNG state in a thread-local variable. You can use `pthread_self()` to seed the RNG differently for all threads.

The use of a mutex is likely to introduce a high overhead. How could the overhead be reduced?

## 2. Implementation of a binary semaphore

Implement your own binary semaphore lock using atomic functions like `atomic_flag_test_and_set()` to make the following program to run:

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <fcntl.h>
4 #include <sys/stat.h>
5 #include <pthread.h>
6
7 int counter = 0;
8 const int nthr = 1000;
9
10 typedef struct {
11     /* MEMBERS TO BE DEFINED */
12 } mysemaphore_t;
13 mysemaphore_t sem;
14
15 int sem_init(mysemaphore_t *s) {
16     /* FUNCTION TO BE IMPLEMENTED */
17     return 0;
18 }
19
20 int sem_wait(mysemaphore_t *s) {
21     /* FUNCTION TO BE IMPLEMENTED */
22     return 0;
23 }
24
25 int sem_post(mysemaphore_t *s) {
26     /* FUNCTION TO BE IMPLEMENTED */
27     return 0;
28 }
29
30 void* func() {
31     sleep(1);
32     sem_wait(&sem);
33     counter++;
34     sem_post(&sem);
35 }
36
37 int main() {
38     pthread_t thr[nthr];
39     sem_init(&sem);
40
41     for (int i = 0; i < nthr; i++)
42         pthread_create(&thr[i], NULL, &func, NULL);
43
44     for (int i = 0; i < nthr; i++)
45         pthread_join(thr[i], NULL);
46
47     printf("counter = %d\n", counter);
48
49     return 0;
50 }
51
52 }
```

Test that the correct results is created using a large number of threads.

## 3. Semaphore with mitigated busy wait

To mitigate the impact of busy waiting, Tom Anderson<sup>1</sup> suggested to add a delay

---

<sup>1</sup>T. Anderson, “The performance of spin lock alternatives for shared-memory multiprocessors”, 1990

after each attempt to acquire a lock. This delay may be incremented in each iteration as shown in the following pseudo-code:

```
type lock = (unlocked, locked)

procedure acquire_lock (L : ^lock)
  delay : integer := 1
  while test_and_set (L) = locked      // returns old value
    pause (delay)                      // consume this many units of time
    delay := delay * 2

procedure release_lock (L : ^lock)
  lock^ := unlocked
```

Update the code implemented in the previous task by adding a delay using the POSIX function `usleep()`. Check whether performance differences can be observed.